Calibrated Hallucination Risk Estimation in Large Language Models Using Evidence-Consistent XGBoost

Md. Atikur Rahaman [a,*], Saiful Alam Sabbir [a], MD. Miraz Ahamed [a], Ohidujjaman Tuhin [a]

[a] *Computer Science and Engineering, United International University, Dhaka 1212, Bangladesh*

ARTICLE INFO

Keywords:
Distribution shift
Evidence-consistent XGBoost
Explainable AI
Hallucination risk estimation
Probability calibration
Question answering
SHAP

Abstract

Large language models answer questions in fluent text, and some answers are wrong in ways the evidence does not support. Detectors for false information have many limitations. For example, they require a large judge model; a computer with graphics capability to run an encoder classifier; feature-based detectors provide little information regarding how well their estimates are calibrated or how well these detectors perform outside the domain of their training data. Therefore, this study was designed to see if a simple black-box model could accurately predict the likelihood of hallucinations from the candidate answer itself along with the question asked and the supporting evidence provided. The proposed method, Evidence Consistent eXtreme Gradient Boosting (EC-XGB), uses an ensemble of decision trees trained over 35 features derived from the input evidence. Eight of the 35 features represent the extent to which the claims are supported. Because the benchmark is nearly fully saturated, the useful gain comes from better reliability during out-of-domain shifts. A tree model using only explicit evidence features requires no model access and remains inexpensive to interpret. Results show that EC-XGB achieves an F1 of 0.9855 and an AUROC of 0.9984 on the HaluEval question answering task, competitive with a conventional XGBoost model. When evaluating results on task type categories (RAGTruth) included in the training set, the AUROC increases from 0.475 to 0.582, while the proportion of natural RAGTruth answers identified as at-risk decreases from 99.9% to 61.3%. Performance on a QA evaluation task not seen during training is consistent with random guessing (i.e., an F1 near 0.30). Finally, a model-specific display calibrator fit to labeled target rows reduces the average target calibration error from 0.726 to 0.126.

1. Introduction

A number of researchers have examined how to apply machine learning techniques to assess the likelihood of a response being a hallucination. One such technique is SelfCheckGPT, which uses several generated answers and assesses their similarity to determine if there has been a hallucination [1]. Another technique is Luna, where a smaller evaluation model is used instead of a larger judge model [2]. A third technique is semantic entropy, which identifies meaning-based groupings of similar responses [3]. A hybrid framework combines frozen encoder features with lightweight classifiers and reports strong benchmark scores [4]. Judges carry high computational costs and encoder models are hard to interpret. Two recent reviews have grouped approaches to detecting hallucinations into three categories including feature-based detectors, encoder-based detectors, and combinations of the above [5, 6]. Both reviews identify evaluation and reliability as two major areas that require additional research, and a third review determines that reliability is an area of concern as well [7].

Six recent studies address the same topic of assessing hallucinations. HaluAgent demonstrated that small language models can select appropriate tools and detect multiple forms of hallucinations, but it required multiple tool calls for each answer [8]. Valentin et al. proposed a calibrated multi-score detector with a detailed cost analysis, but the detector relies heavily on large language model (LLM) scoring for its primary source of information [9]. Sundaragiri et al. proposed a system that uses retrieval, natural language inference (NLI) verification, and an XGBoost classifier, and they report only classification metrics [10]. Haq et al. assigned SHAP-LIME attributions to each word, but did not test whether those values remained consistent across words [11]. Deng et al. found that detectors degrade once they leave the domain in which they were trained, and their method focuses on multi-turn dialogue signals [12]. FaithBench shows that modern LLMs still produce hard summarization hallucinations, and it does not assess calibration risks [13].

*Corresponding author: Md. Atikur Rahaman
Email addresses: mrahaman2310298@bscse.uiu.ac.bd (Md.A. Rahaman); sabbir2310105@bscse.uiu.ac.bd (S.A. Sabbir); mahamed2310524@bscse.uiu.ac.bd (MD.M. Ahamed)

The existing literature leaves four issues unaddressed. Calibration is infrequently reported for lightweight detectors, explanation quality is rarely measured, cross-domain behavior is rarely tested, and error patterns are rarely analyzed case by case. This paper addresses all four issues by developing a black-box model that takes a query , $are f erencecontext$, and a candidate answer , $andoutputs(H \mid Q, C, A)$. The model, evidence-consistent XGBoost (EC-XGB), consists of 26 evidence-consistency features in seven groups, eight claim-level NLI aggregates, and a tuned XGBoost classifier. Platt and isotonic calibration yield the probability, SHAP explanations reveal feature contributions, artifact-control baselines and group-aware splits prevent shortcut solutions, and validation-only calibration solves the reliability problem. Transfer limits are established on RAGTruth and FaithBench.

HaluEval QA provides 10,000 paired correct and hallucinated answers for training and in-domain testing [14]. RAGTruth provides natural responses for transfer and recalibration [15]. The scope is question answering with a reference context, and FaithBench serves only as an out-of-domain summarization stress test. EC-XGB reaches an F1 of 0.9855 on HaluEval QA. On natural RAGTruth answers, the flag rate falls from 99.9% to 61.3%, and the AUROC rises from 0.475 to 0.582. Target calibration lowers the RAGTruth QA ECE from 0.726 to 0.126. On HaluEval, recall stays above 99% at a five percent false-positive rate, and one analysis takes 61.8 ms at about 0̇01 per 1,000 predictions.

2. Conventional Method

2.1. Preliminaries

Fluent hallucinations are those that are not supported by the available evidence. In most cases, detection quality is assessed using labeled benchmarks and a metric, typically precision, recall, F1 or the area under the receiver operating characteristic curve (AUROC). Current test beds include HaluEval [14] and TruthfulQA [16] for hallucination detection. FActScore [17] assesses factual accuracy for long-form generation. For deployment, there are two key evaluation considerations. First, calibration. Only if a risk score has a numeric value that is consistent with the probability of generating a hallucination can that risk score be used effectively. Expected calibration error (ECE) and reliability diagrams are used to show how far off a model is in terms of calibration [18]. Second, the need for interpretability. Users need to understand what specific signal caused their high-risk flag to appear, and this can be done through feature attribution techniques such as SHapley Additive exPlanations (SHAP) for tabular models [19]. This work assesses both of these properties within a black-box setting.

2.2. Baseline Method

The conventional method in this study is a tuned XGBoost classifier [20] over the 26 hand-crafted evidence-consistency features. XGBoost is selected as the baseline because it handles nonlinear interactions, it trains quickly on a tabular design matrix, and its tree structure supports exact SHAP values. XGBoost is a gradient-boosting ensemble, so the model adds decision trees one at a time and each new tree fits the error left by the trees already in the ensemble [20]. Each tree splits the feature space into leaves, and each leaf stores a weight that adds to the running score. The final score passes through a logistic link, so the output is a probability of the hallucinated class.

Let x_i be the feature vector of row i and $y_i \in \{0, 1\}$ its label. The model sums K trees, $z_i = \sum_{k=1}^{K} f_k(x_i)$. The training objective combines a differentiable loss with a regularizer over the trees,

$$\mathcal{L}(\phi) = \sum_{i=1}^{n} l(y_i, \hat{y}_i) + \sum_{k=1}^{K} \left(\gamma T_k + \frac{1}{2}\lambda \|w_k\|^2\right), \tag{1}$$

where $\phi = \{f_k\}_{k=1}^{K}$ collects the trees, f_k is the k-th tree with T_k leaves of weight w_k, γ is a leaf-count penalty, and λ is an L2 weight penalty. The loss l is the standard logistic loss and the probability comes from the sigmoid of the summed leaf score, $\hat{y}_i = \sigma(z_i) = 1/(1 + e^{-z_i})$.

Tree depth, learning rate, tree count, subsample rate, and column subsample rate are tuned on the training split, and the selected values and the search protocol are reported in Section 4. This baseline is the reference for every later modification, and Algorithm 1 summarizes its training and inference procedure.

Algorithm 1 Baseline feature-based XGBoost (training and inference)

Require: labeled rows (q_i, c_i, a_i, y_i) and seeds $S = \{42, 123, 456\}$
Ensure: trained model, calibrator, and decision threshold
1: extract the 26 evidence-consistency features for every row
2: tune XGBoost with a randomized search and grouped 5-fold cross-validation
3: **for** each seed $s \in S$ **do**
4: train the classifier on the training split
5: fit Platt calibration and select the decision threshold on the validation split
6: evaluate on the fixed test split
7: **end for**
8: **return** the mean metrics over the seeds

2.3. Limitations and Research Gap

This baseline performs well as a standalone for its specific benchmark, but it is limited. On HaluEval the benchmark has saturated, so the addition of features has little impact on performance. The baseline relies heavily on both answer length and lexical similarity. As demonstrated in Section 5, answer length is essentially an ideal feature to utilize when working with this dataset. However, under domain shift, the raw probabilities lose their reliability. Specifically, the baseline flags virtually all natural answers, and its confidence no longer correlates with the observed error rate.

The benchmark itself represents one limiting factor, since close to 85% of HaluEval answers were produced by sampling and filtering [21]. The line separating factual errors from hallucinations is therefore blurred in many instances. Another issue is that the baseline does not include a claim-level evidence signal, was trained on a single corpus, and provides a probability without a method to calibrate it to the target domain or a reliability test for its explanations.

3. Proposed Method

3.1. Working Principle and Task

Figure 1 shows the full proposed pipeline at a glance. The system takes a question, an optional reference context, and a candidate answer produced by a black-box LLM, and it returns a calibrated risk probability together with per-feature explanations.

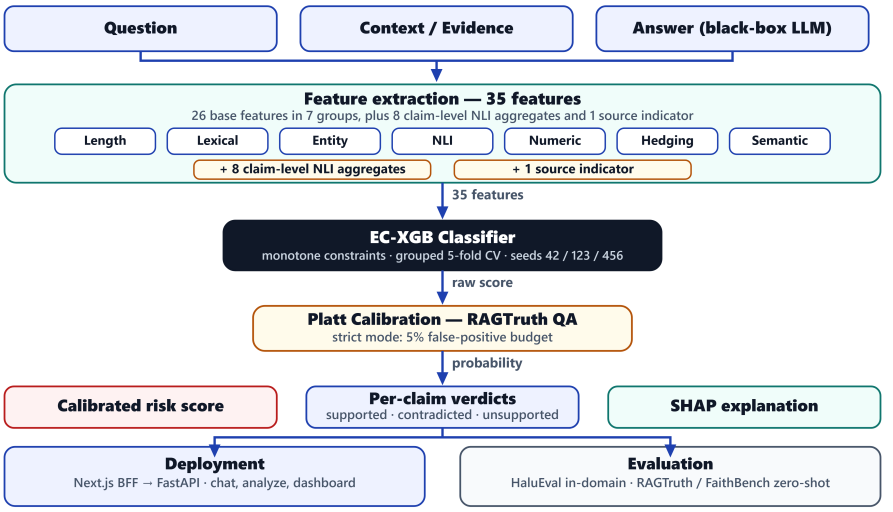


Figure 1: End-to-end architecture of the proposed model. The inputs feed 35 features into the EC-XGB classifier, Platt calibration produces the deployed score, and the outputs are the risk score, the per-claim verdicts, and the SHAP explanations. The system is deployed as a Next.js and FastAPI application and evaluated in-domain on HaluEval and under shift on RAGTruth and FaithBench.

Given a question q, an optional reference context c (evidence), and a candidate answer a produced by a black-box LLM, the task is to predict the probability $P(\text{hallucinated} \mid q, c, a)$. The proposed model predicts risk, not truth. It flags answers that are unsupported by, or contradictory to, the available evidence.

3.2. Features

Table 1 lists the 26 base features in seven groups, and Table 2 lists the eight claim-level aggregates that the proposed model adds. Natural language inference (NLI) decides whether one sentence follows from another. The NLI signals come from a DeBERTa-v3 cross-encoder [22] trained on SNLI [23] and MultiNLI [24], evaluated in both directions, context-to-answer and answer-to-context, which yields six probability features. Semantic similarity uses MiniLM sentence embeddings [25, 26], entity features use spaCy NER [27], and the numeric and hedging groups are collected with regular expressions and a small hedging lexicon.

Table 1
Base feature groups (26 features total).

Group	Features	Source
Length/style	n_chars, n_words, n_sentences, avg_word_len	regex
Lexical	overlap vs. context & question, Jaccard ×2	token sets
Entity	entity counts, overlap ratio, novel-entity ratio	spaCy NER
NLI	entail/contradict/neutral × 2 directions	cross-encoder NLI
Numeric	number counts, overlap ratio, novel numbers	regex
Hedging	hedge count, hedge density	lexicon
Semantic	cosine(context,answer), cosine(question,answer)	MiniLM

Table 2
Claim-level NLI aggregates (8 features). The answer is split into at most six atomic clauses, with conjunctions such as "and" or "but" broken into separate clauses, each clause is scored against its four best-matching context sentences, and the pair probabilities are aggregated.

Signal	Features	Source
Contradiction	max and mean contradiction, contradicted-claim ratio	NLI cross-encoder
Support	supported-claim ratio, unsupported-claim ratio, min and mean entailment	NLI cross-encoder
Decomposition	number of claims	sentence splitter

3.3. Evidence-Consistent XGBoost

EC-XGB keeps the base XGBoost learner and changes three things, added cumulatively so that each change can be measured on its own. The first change replaces the raw length counts with log-scaled values and adds monotone constraints, so risk cannot fall when contradiction or novelty rises and cannot climb when entailment, overlap, or semantic similarity rises. A log-scaled length feature replaces the raw count, $\ell' = \log(1 + \ell)$, which compresses the long tail of the length distribution. A monotone constraint fixes the sign of the model response for a chosen feature, so an evidence-positive feature x_j keeps $\partial\hat{y}/\partial x_j \leq 0$, and the sign flips for an evidence-negative feature such as contradiction. In plain terms, more contradiction can never lower the risk score, and more entailment can never raise it.

The second change appends the eight claim-level aggregates from Table 2, which lets the classifier see the same evidence signals that the deployed verification layer reports. Take an answer split into K claims, with c_k the contradiction and e_k the entailment of claim k. The aggregate vector and the full model input are

$$\psi(x) = \Big[\max_k c_k,\ \bar{c},\ r_{\text{con}},\ r_{\text{sup}},\ r_{\text{unsup}},\ \min_k e_k,\ \bar{e},\ K\Big], \tag{2}$$

$$\phi'(x) = \left[\ell',\ \tilde{\phi}(x),\ \psi(x),\ s\right], \tag{3}$$

where r_{con}, r_{sup}, and r_{unsup} are the contradicted, supported, and unsupported claim ratios, and $\tilde{\phi}(x)$ holds the remaining base features after the length change, which gives 35 inputs in total.

The third change adds RAGTruth summarization and data-to-text training rows with a binary source indicator $s \in \{0, 1\}$, while every RAGTruth QA row stays out of training, which keeps the held-out QA classifier zero-shot. The deployed display score is fitted on labeled target rows and is not zero-shot. The deployed checkpoint combines all three changes and is called EC-XGB.

Three protocols are distinguished. P0 is source-only: the calibrator and the threshold are fitted on HaluEval validation. P1 is target-calibrated: the calibrator is fitted on the 5,034 RAGTruth QA calibration rows and the threshold is fixed at 0.5. P2 is target-thresholded: the threshold is selected on the same target calibration split at a five percent false-positive budget. Algorithm 2 lists the full procedure.

Algorithm 2 Evidence-consistent XGBoost (EC-XGB) training and deployment

Require: labeled HaluEval QA rows, RAGTruth non-QA rows, seeds S
Ensure: deployed model, display calibrator, source threshold, target threshold
1: extract the 26 base features and the eight claim-level aggregates
2: log-scale the length features and set the monotone constraints
3: **for** each seed $s \in S$ **do**
4: build the 35-feature input with the source indicator
5: train XGBoost on HaluEval plus RAGTruth non-QA rows, and keep every RAGTruth QA row out of training
6: fit the display calibrator on the RAGTruth QA calibration split (P1)
7: select the source-only threshold on the HaluEval validation split (P0)
8: select the target threshold on the same calibration split at five percent FPR (P2)
9: **end for**
10: **return** the deployed model, calibrator, and thresholds

The contribution is the combination. Each piece is standard on its own, and together they hold the in-domain score, cut the false-alarm rate under shift, and keep the output explainable at the feature level. The contribution is therefore a method and an evaluation protocol rather than a new learning algorithm. The deployed system adds claim-level verification with evidence retrieval and a display calibrator fitted on the target domain. Relative to the closest feature-based detector [10], the difference is the claim-level evidence features and the calibration protocol. Relative to the LLM-scored detector [9] and the tool-based detector [8], the difference is that inference needs no model calls. Their scores come from other corpora and other protocols, so the comparison is about the protocol and not about the numbers. None of the three reports the calibration error or the shifted flag rate that this study measures.

3.4. Calibration and Explainability

Calibrators are fit on the validation split only. Platt scaling fits a sigmoid $p(x) = \sigma(Af(x) + B)$ to the raw score $f(x)$, where A and B are fit by maximum likelihood [28]. Isotonic regression fits a non-decreasing step function g so that $p(x) = g(f(x))$ [29]. Both are compared against the raw model on the test split. The comparison uses expected calibration error over $M = 10$ bins,

$$\text{ECE} = \sum_{m=1}^{M} \frac{|B_m|}{n} \left|\text{acc}(B_m) - \text{conf}(B_m)\right|, \tag{4}$$

where $\text{acc}(B_m)$ and $\text{conf}(B_m)$ are the accuracy and mean confidence in bin B_m, together with the Brier score $\text{BS} = \frac{1}{n}\sum_i (\hat{p}_i - y_i)^2$ and reliability diagrams [18]. A second step targets a new domain. The calibrator can be re-fit on a small labeled sample of the target corpus, and the effect is measured on a disjoint held-out part of the target corpus. The deployed EC-XGB display score follows the same rule: its calibrator is fitted on the 5,034-row RAGTruth QA calibration split, and the source indicator is set to one for natural responses. This target calibration needs labeled target rows and is not zero-shot. The source-only protocol uses the HaluEval validation split for both the calibrator and the threshold.

Global explanations come from SHAP [19]: a TreeExplainer over the raw XGBoost model produces mean absolute importance and beeswarm summaries, and local explanations show the feature contributions behind a single verdict. Explanation reliability is measured directly in this work, since feature-attribution methods can be unstable [30]. Each

test answer is perturbed in six ways, by replacing entities, numbers, or dates, removing supporting context, inserting irrelevant context, and shuffling clauses. Stability is quantified with Spearman and Kendall rank correlation between the original and perturbed scores, the fraction of top-1 flips, and the top-k Jaccard overlap of the SHAP feature rankings.

3.5. System and Reproducibility

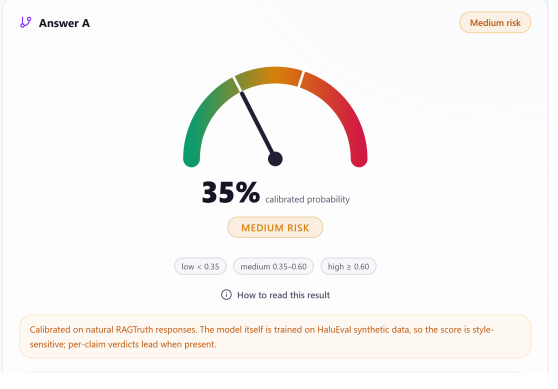
The methods above are deployed as a two-tier web application. The Next.js frontend streams an LLM conversation on the chat page using the assistant-ui library, and the FastAPI backend holds the model, the calibrator, the SHAP explainer, and the feature models, all loaded once at startup. A Next.js proxy forwards every request to the FastAPI endpoints, which avoids cross-origin issues and keeps the API key on the server. All quantitative results in Section 5 measure the EC-XGB classifier and its calibrator alone. The claim-verification layer described below is a separate deployment component. It changes the displayed verdict and is demonstrated through the interface rather than measured in the reported metrics.

The chat page streams through the Vercel AI SDK route. A tool call from the conversation model triggers the analysis endpoint, and the result renders as a risk card with the calibrated score, the per-claim verdicts, the SHAP bars, and an evidence panel.

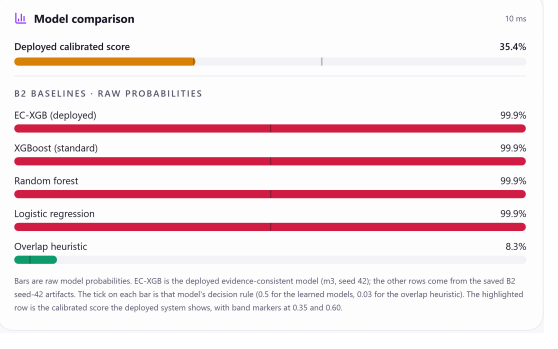
Beyond the score, the system verifies the answer claim by claim. Each claim is decomposed from the answer text and checked with sentence-level NLI against one of three evidence sources, the provided context, an indexed document collection, or live web search. Verdicts are decided support-first, so any entailing sentence supports a claim and a contradiction decides only when no sentence supports it. A relevance gate then drops retrieved sentences that do not cover the claim vocabulary. The verdicts supported, contradicted, and unsupported are rendered per claim, with the driving evidence sentence quoted for contradicted claims. Web search uses the Brave Search API, and the document index combines BM25, FAISS dense retrieval, and reciprocal rank fusion with a cross-encoder reranker. An uncertain verdict routes the claim to an LLM judge.

The verdict headline leads the card. Any contradicted claim sets the risk to high, an unsupported claim sets it to medium, and fully supported claims set it to low. This ordering follows the style-sensitivity result in Section 5, where a full-sentence grounded answer can saturate the calibrated score even though every claim is supported. The displayed percentage comes from the EC-XGB display calibrator and is adjusted by the claim verdicts.

Figure 2 shows the analyze page for the hallucinated-number case, where the context states 5 days per decade and the answer states 10. The raw EC-XGB probability is 0.999, so the raw score alone is not an actionable risk level. The deployed display calibrator reports 35.4% and marks the answer as medium risk. The comparison panel shows the same pattern across the baselines, with the learned models all near 99.9% and the overlap heuristic at 8.3%, which is why the interface leads with the calibrated score and keeps the raw score as a secondary signal.



(a) Risk card: display score 35.4% (medium risk).



(b) Model comparison: raw probabilities of the deployed model and the baselines.

Figure 2: Analyze page for the hallucinated-number case, where the context states 5 days per decade and the answer states 10. The risk card shows the calibrated display score, and the comparison panel shows the raw probability of each model. The raw EC-XGB probability is 0.999, while the deployed display calibrator reports 35.4%, which makes the risk level actionable.

The source repository contains the complete pipeline: `src/data` holds the download and split scripts, `src/features` the seven feature groups, `src/models` the training protocol with the cross-domain, calibration-shift, and explanation-reliability runs, `src/explain` the SHAP analysis, and `src/api` the FastAPI application. `requirements.txt` pins every dependency with exact versions, split indices are saved to `artifacts/split_indices.json` and `.npy`, and all artifacts live under `artifacts/models/`. A self-contained Colab notebook reproduces the run on a GPU, `run_all_experiments.py` reproduces it locally from `configs/version_b.yaml`, and `verify_artifacts.py` checks every reported number against the frozen artifacts.

4. Experiments

4.1. Experimental Setup and Conditions

All experiments run on a Windows 11 laptop with an NVIDIA RTX 3060 6 GB GPU (CUDA 12.8, torch 2.11.0+cu128) and 32 GB of RAM, and the code runs on Python 3.12. The model uses scikit-learn 1.9.0, XGBoost 3.3.0, SHAP 0.52.0, spaCy 3.8.14, and sentence-transformers 5.6.1. Saved split indices and centralized seeds make every run reproducible.

4.2. Dataset

The QA subset of HaluEval [14] is the training source. It contains 10,000 questions, each with a knowledge passage, a correct answer, and a hallucinated answer, which gives 20,000 labeled rows. The rows are split 70/15/15 at the question-group level, producing 14,000 training, 3,000 validation, and 3,000 test rows. Both answer variants of a question share the same context, so a random split would let the model memorize the pairing.

Transfer is measured on RAGTruth [15], a corpus of naturally generated responses with human word-level span annotations, and on FaithBench [13], a summarization hallucination benchmark built from modern LLM outputs. The RAGTruth QA task contributes 5,934 rows in 989 groups, of which 5,034 rows (839 groups) form the recalibration set and 900 rows (150 groups) the held-out test, with no overlapping groups between the two. Table 3 summarizes the corpora.

Table 3
Datasets used for training and evaluation.

Corpus	Rows	Groups	Use
HaluEval QA	20,000	10,000	train / validation / test
RAGTruth QA	5,934	989	held-out + recalibration
RAGTruth full	17,790	2,965	transfer (seen tasks)
FaithBench	750	750	zero-shot

Text is whitespace-normalized. Lexical features operate on lowercased token sets, while NER and NLI keep the original casing. Missing context is handled explicitly, with lexical overlap set to zero and NLI probabilities to one third. All three corpora are used under their original licenses: HaluEval and RAGTruth are MIT-licensed, and FaithBench is distributed under CC BY-NC-SA 4.0. Raw files are not redistributed. The repository ships the download scripts, source revisions, and SHA-256 hashes instead, so every split can be regenerated from the original repositories.

4.3. Experimental Procedure

Nine models are compared. The overlap heuristic sets risk to one minus the lexical overlap between the answer and the context, and its threshold is tuned on validation. Logistic regression on standardized features and a random forest with 300 trees are the two standard classifiers. The conventional method is XGBoost [20], trained on the 26 base features, and it is the reference for the modified model in Section 3. Three controls isolate the value of the feature design, an NLI-only logistic regression that sees only the six NLI probabilities, an answer-only TF-IDF linear model, and a TF-IDF linear model over the full input. Internal-state detectors read hidden activations and need model access, so they sit outside the black-box setting [31]. A majority baseline that predicts the negative class for every row confirms the balance of the labels.

A fine-tuned RoBERTa-base classifier gives the source-only neural comparison [32]. It reads the answer with its question and context as one 256-token text on the same grouped split. Training runs for three epochs with binary

cross-entropy, and Platt calibration is fitted on validation, with only 14 of 3,000 test rows exceeding 256 tokens. One seed takes about 15 minutes on an RTX 3060 6 GB card.

4.4. Parameter Configuration

Hyperparameters come from a randomized search over 30 configurations, scored with grouped 5-fold cross-validation on the training split. For XGBoost the search selects depth 4, learning rate 0.01, 500 trees, subsample 0.9, and column subsample 0.7, with a cross-validated AUROC of 0.9970. The neural baseline uses 256-token inputs, three epochs, and binary cross-entropy, and its Platt calibrator is fitted on validation, never on test rows. Table 4 fixes the calibrator and threshold for each protocol, and seeds 42, 123, and 456 are repeated throughout.

4.5. Evaluation Metrics

Evaluation uses precision, recall, F1, AUROC, the area under the precision-recall curve (PR-AUC), and the Matthews correlation coefficient (MCC). Calibration is measured with expected calibration error over 10 bins, the Brier score, and negative log-likelihood. Metric tables report four decimals, and latency is reported in milliseconds. Significance is assessed with McNemar's test on the paired seed-42 test predictions, bootstrap 95% confidence intervals from 1,000 resamples over question or source groups, and a Wilcoxon signed-rank test across seeds. The group structure is respected everywhere, so the two rows of one question never appear in both the training and the evaluation folds.

4.6. Additional Protocols

Raw, Platt, and isotonic probabilities are compared on the test split. Calibration is known to degrade under distribution shift [33], so the target-domain protocol follows the same principle. Section 3 names the three protocols (P0, P1, and P2, summarized in Table 4) and the split that fits each one. All P1 and P2 numbers come from the disjoint 900-row test set, and no test label enters either fit.

Table 4
The three evaluation protocols. Only P0 is zero-shot on the target corpus, because P1 and P2 both read labels from the target calibration split. The 900-row test split is disjoint from that calibration split in every row.

Protocol	Calibrator fit	Threshold fit	Zero-shot
P0 source-only	HaluEval validation (3,000)	HaluEval validation	yes
P1 target-calibrated	RAGTruth QA calibration (5,034)	fixed at 0.5	no
P2 target-thresholded	RAGTruth QA calibration (5,034)	same rows at 5% FPR	no

Global explanations come from a SHAP TreeExplainer [19] over the raw XGBoost model. Three measurements check whether the explanations are trustworthy. First, the mean absolute SHAP ranking is compared against permutation importance with a Kendall tau correlation. Second, the top-5 feature set is re-estimated from 1,000 bootstrap resamples and compared with the Jaccard index.

Third, two perturbation studies are run on the test answers. Content edits replace entities, numbers, or dates, remove the supporting context, or insert an irrelevant sentence, and each edit should move the calibrated score if it changes the evidence. Structure edits shuffle the clause order, and these should leave the score nearly unchanged. For each operator the mean absolute score change, the fraction of large score changes, the top-1 SHAP flip rate, and the Spearman correlation between the original and perturbed scores are recorded. A neutralization study sets the top-k SHAP features of each sample to its median and tracks the mean score change for $k = 1, 3, 5, 10$.

Latency is measured on 200 test samples on the reference laptop with an RTX 3060 6 GB GPU, where NLI and embeddings run on the GPU and the remaining features on the CPU. The LLM-as-judge comparison uses GPT 5.6 Luna, a standard but biased comparator [34], on the same 200 samples with a fixed prompt, and the cost is computed from the recorded token counts.

5. Results

5.1. Model comparison

Table 5 reports the test-set performance of every baseline, averaged over the three seeds. XGBoost reaches an F1 of 0.9846 with an AUROC of 0.9979, ahead of the random forest by 0.0004 in F1 and the logistic regression by 0.0186. The overlap heuristic, which needs no training beyond its threshold, still reaches an F1 of 0.9352, and a majority

baseline scores zero, which confirms balanced labels. The NLI-only baseline lags at 0.6714, which shows that the lexical, length, and semantic signals carry information the NLI cross-encoder alone does not. The TF-IDF variants sit between these extremes, with an answer-only linear model at 0.9224 and a full-input model at 0.5999.

The XGBoost margin over the random forest is small but significant in a McNemar test on the paired seed-42 test predictions ($p = 0.044$). The logistic regression ($p = 0.43$) and heuristic ($p = 0.35$) comparisons are not. Bootstrap confidence intervals over 1,000 question-group resamples for XGBoost are F1 [0.9809, 0.9896] and AUROC [0.9964, 0.9989]. The conventional XGBoost row in Table 5 and the m0 row in Table 6 come from the baseline and modified pipelines.

The source-only neural baseline wins in-domain (F1 0.9912 against 0.9846) but collapses under shift, with a RAGTruth AUROC of 0.337 against 0.582 for EC-XGB. An adapted encoder trained on the same multi-source rows as EC-XGB closes the seen-task gap and passes EC-XGB there (0.694 against 0.582). It does not transfer: on the held-out QA task it stays near chance (0.490 against 0.549), and on FaithBench it falls behind (0.434 against 0.594). Adaptation lifts any model on seen tasks, while the feature model keeps more of its ranking on unseen ones.

The split design is checked against a known confound. A row-level split, where the two rows of one question can fall into different partitions, inflates F1 to 0.9886. The grouped split corrects this to 0.9846, a drop of 0.004 in F1 and 0.0001 in AUROC.

Table 5
Test-set performance on HaluEval QA, mean over seeds 42/123/456 (n = 3,000). The ECE and Brier columns measure the reliability of each model's risk score. The neural baseline is a fine-tuned RoBERTa-base with 256-token inputs. The majority baseline predicts the negative class for every row. Probability-based metrics are undefined for a constant score.

Model	P	R	F1	AUROC	PR-AUC	MCC	ECE	Brier
Heuristic (1−overlap)	0.9473	0.9233	0.9352	0.9132	0.8196	0.8723	0.3541	0.2672
Logistic regression	0.9768	0.9553	0.9660	0.9932	0.9901	0.9329	0.0099	0.0253
Random forest	0.9897	0.9789	0.9842	0.9980	0.9983	0.9687	0.0143	0.0127
NLI-only	0.6479	0.6967	0.6714	0.7177	0.7531	0.3189	0.0678	0.2182
TF-IDF (answer-only)	0.9550	0.8920	0.9224	0.9696	0.9665	0.8519	0.0831	0.0662
TF-IDF (full input)	0.6397	0.5647	0.5999	0.6754	0.6785	0.2484	0.0679	0.2328
XGBoost (conventional)	0.9932	0.9762	0.9846	0.9979	0.9983	0.9697	0.0045	0.0117
RoBERTa-base	0.9955	0.9869	0.9912	0.9990	0.9993	0.9825	0.0070	0.0081
Majority (negative)	0.0000	0.0000	0.0000	—	—	0.0000	—	—

Single-answer behavior shows where aggregate metrics hide disagreements. For a grounded span and a hallucinated sentence that share one question and one context, the overlap heuristic and the NLI-only model fail to separate the pair. Logistic regression, random forest, and XGBoost separate it cleanly.

5.2. Evidence-consistent XGBoost

Table 6 reports the component ablation for EC-XGB on the HaluEval test set, averaged over the three seeds. The conventional XGBoost reference (m0) reaches an F1 of 0.9844, and the length debias with monotone constraints (m1) reaches 0.9852. The claim-level aggregates (m2) reach 0.9857 (+0.0005 over m1, within seed-level variation), and the full EC-XGB (m3) reaches 0.9855 with an MCC of 0.9715. The differences are small and not significant in a McNemar test on the paired predictions (m0 against m3, $p = 0.45$), and a Wilcoxon test across seeds finds no difference either ($p = 0.25$).

The benchmark is saturated, so this ablation is a preservation check rather than a win. Every component holds the line. Removing only the monotone constraints leaves the F1 unchanged, lowers the ECE by 0.0009, and raises the RAGTruth AUROC by 0.0043. On the seed-42 split false positives fall from 12 to 8. Removing the source indicator from the same recipe changes the in-domain F1 by less than 0.001 (0.9849 against 0.9855). The shifted flag rate moves by less than one point, so the multi-source gain does not rest on source identity. At inference, the indicator is set to one for natural responses and zero for benchmark-style rows.

Table 7 shows where the modification matters. Every model is applied without retraining to the RAGTruth and FaithBench corpora. The conventional model flags almost every natural answer as risky, since 99.9% of RAGTruth rows and 99.5% of FaithBench rows cross the 0.5 threshold, which makes the decision rule useless outside the benchmark even though recall stays high. EC-XGB cuts the flag rate (the share of answers above the 0.5 threshold) to 61.3% on the RAGTruth test tasks and to 7.3% on FaithBench. It raises the AUROC from 0.475 to 0.582 and from 0.509 to 0.594,

Table 6
EC-XGB component ablation on the HaluEval test set, mean over seeds 42/123/456 (n = 3,000). m0 is the conventional XGBoost reference, m1 adds log-scaled length features and monotone constraints, m2 adds the eight claim-level aggregates (34 features), and m3 adds the binary source indicator (35 features) and is the deployed EC-XGB.

Variant	Features	F1	AUROC	MCC	ECE
m0 conventional	26	0.9844	0.9980	0.9693	0.0066
m1 debias	26	0.9852	0.9976	0.9708	0.0051
m2 + claims	34	0.9857	0.9984	0.9717	0.0078
m3 EC-XGB	35	0.9855	0.9984	0.9715	0.0074

and it lowers the ECE on RAGTruth from 0.635 to 0.278. The held-out RAGTruth QA task is the strictest test, since its rows never enter training. There the F1 stays flat (0.3022 against 0.3081) and the ECE improves from 0.815 to 0.743, which shows that multi-source training does not transfer a new task for free. On FaithBench the conservative regime lowers the F1 from 0.813 to 0.120. The conventional model earns that F1 by flagging nearly every row, so the AUROC and the flag rate are the fair comparison at this operating point, and the F1 drop is reported as a limitation.

Table 7
Cross-domain behavior, mean over seeds. Conventional is the B2 XGBoost trained on HaluEval only. Its external results are zero-shot. EC-XGB adds the RAGTruth non-QA training rows, so its summarization and data-to-text results measure adaptation, while its QA results are held out. RAGTruth all uses the 2,700-row test split, RAGTruth QA uses the 900-row held-out split, and FaithBench uses all 750 rows. Flagged is the share of rows above the 0.5 threshold.

Corpus	Model	F1	AUROC	Flagged	ECE
RAGTruth all	Conventional	0.5181	0.4753	99.9%	0.6346
RAGTruth all	EC-XGB	0.5608	0.5815	61.3%	0.2782
RAGTruth QA	Conventional	0.3022	0.5465	99.9%	0.8149
RAGTruth QA	EC-XGB	0.3081	0.5491	97.4%	0.7433
FaithBench	Conventional	0.8130	0.5089	99.5%	0.2917
FaithBench	EC-XGB	0.1201	0.5941	7.3%	0.4500

The source-only operating point (P0) selects the threshold on the HaluEval validation split that maximizes recall under a five percent false-positive budget. The test false-positive rate stays near 5% while recall stays above 99%, a recall-first option at about one point of F1. Table 9 puts both models on the same 900-row RAGTruth QA test set. Source calibration barely moves the ECE, while target Platt cuts it to 0.1335 for the conventional model and 0.1263 for EC-XGB, with isotonic close behind. The target-calibrated display score (P1) is the seed-42 target Platt fit, which moves the raw ECE from 0.726 to 0.126. The table reports three-seed means, which is why its raw and calibrated cells read 0.7433 and 0.1263.

AUROC is unchanged by calibration, and every target-calibrated score falls below 0.5, so the fixed threshold predicts no positives. The target-thresholded mode (P2) selects a threshold on the target calibration split at a five percent false-positive budget (0.313). It recovers a test recall of 5.0% at a 5.8% false-positive rate (precision 15.7%, F1 0.076; Wilson 95% intervals: recall [2.6, 9.6], FPR [4.3, 7.7]). The test rate sits above the 5% target because the threshold is fixed on calibration. Target recalibration fixes the probability scale, not the ranking. Already 252 labeled target rows give a test ECE of 0.133.

5.3. Calibration

Table 8 compares raw, Platt-scaled, and isotonic probabilities: the raw model is already well calibrated (ECE 0.0045), neither recalibration improves it, and the reliability diagrams confirm the same.

Figure 4 shows the raw and target-calibrated score distributions behind the conventional-model rows. The practical constraint is the label budget: 252 labeled target rows already give the same test ECE as the full 5,034.

5.4. Cross-domain transfer

The transfer evaluation is the central negative result. The conventional model is applied without adaptation, so its external numbers are zero-shot. EC-XGB has seen the RAGTruth summarization and data-to-text rows, so those results

Table 8
Calibration on the HaluEval test set (mean over seeds).

Method	F1	ECE	Brier	NLL
Raw XGBoost	0.9846	0.0045	0.0117	0.0449
Platt (sigmoid)	0.9848	0.0091	0.0129	0.0579
Isotonic	0.9844	0.0070	0.0124	0.0727

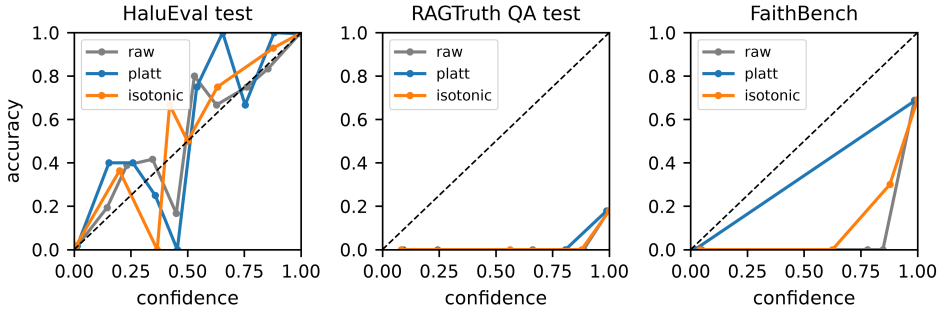


Figure 3: Reliability diagrams of the raw, Platt-calibrated, and isotonic models on the HaluEval test set (left), RAGTruth QA (middle), and FaithBench (right).

Table 9
Calibration on the same 900-row RAGTruth QA test set (17.8% positive; means over seeds 42/123/456). The deployed EC-XGB display calibrator is the seed-42 target Platt row (ECE 0.126). AUROC is unchanged by calibration: 0.540 for the conventional model and 0.549 for EC-XGB.

Model	Calibrator	ECE	Brier	NLL	Pred. pos.
Conventional	Raw	0.8185	0.8164	5.20	0.999
Conventional	Source Platt (HaluEval)	0.8089	0.8011	3.62	0.999
Conventional	Target Platt (RAGTruth)	0.1335	0.1639	0.51	0.000
Conventional	Target isotonic	0.1316	0.1663	0.53	0.000
EC-XGB	Raw	0.7433	0.7061	2.91	0.974
EC-XGB	Source Platt (HaluEval)	0.7735	0.7527	3.18	0.975
EC-XGB	Target Platt (RAGTruth)	0.1263	0.1613	0.50	0.000
EC-XGB	Target isotonic	0.1260	0.1618	0.50	0.000

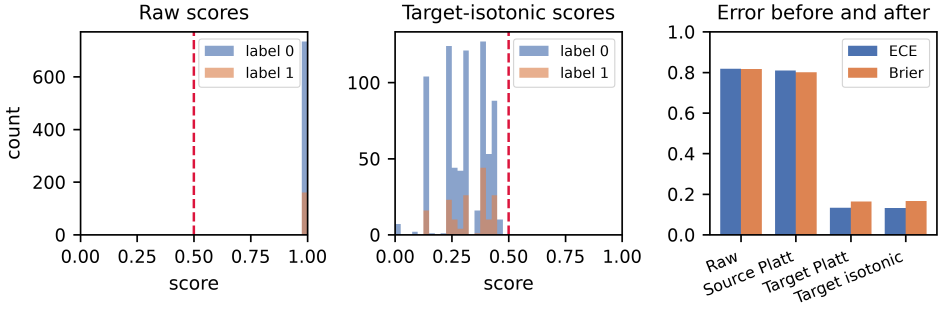


Figure 4: Target-domain recalibration of the conventional XGBoost model on RAGTruth QA. Left and middle: raw and target-calibrated score distributions on the 900-row test set, with the 0.5 threshold marked. Right: ECE and Brier before and after recalibration.

measure adaptation. The EC-XGB held-out QA classifier is zero-shot, but the QA display score is target-calibrated and is not. FaithBench is scored with the source calibrator because no official calibration split exists.

On the full 17,790-row RAGTruth corpus the conventional model reaches an F1 of 0.603, an AUROC of 0.497, and an ECE of 0.560. Its recall is 1.0 on every corpus: it flags almost everything, precision lands on the label prior, and the AUROC stays near 0.5 because synthetic HaluEval hallucinations differ from natural responses. On FaithBench the same model reaches an F1 of 0.8130 and an AUROC of 0.5314. Task populations repeat the pattern (F1 0.4515 QA, 0.4603 summarization, 0.8140 data-to-text), rising from 0.4940 on very short contexts to 0.7091 at 128 to 511 characters. Table 7 reports the EC-XGB comparison on the 2,700-row test split, whose lower positive rate (35% against 43%) explains the lower level. Group-aware bootstrap intervals (1,000 resamples) keep the held-out QA AUROC near chance ([0.514, 0.607] conventional, [0.502, 0.603] EC-XGB), and Figure 5 plots the score distributions.

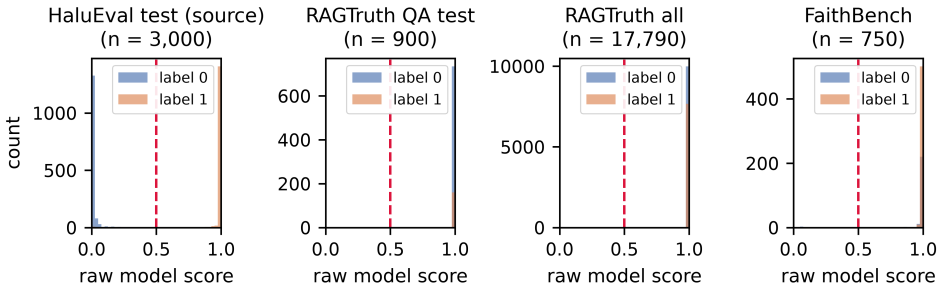


Figure 5: Raw model score distributions on the HaluEval test set and the external corpora, with the 0.5 threshold marked. The source panel separates the classes, while every external panel piles both labels at 1.0.

5.5. Ablation

Table 10 removes each feature group and retrains from scratch. Lexical features dominate (−0.0247 F1), length and NLI contribute less, entity, numeric, and hedging stay within 0.0002 F1 of the full model, and removing the semantic group slightly improves the F1, and the paper keeps that negative result.

Table 10
Feature-group ablation, mean over seeds (full model F1 = 0.9846, AUROC = 0.9979).

Removed group	F1	ΔF1	AUROC
length	0.9801	−0.0045	0.9963
lexical	0.9600	−0.0247	0.9923
entity	0.9847	+0.0001	0.9982
nli	0.9823	−0.0023	0.9973
numeric	0.9844	−0.0002	0.9982
hedging	0.9848	+0.0001	0.9981
semantic	0.9860	+0.0014	0.9982

5.6. Explanation reliability

Figure 6 shows the mean absolute SHAP values for the top-12 features. The context overlap dominates, followed by the answer length features, the context-to-answer contradiction probability, and the question-to-answer cosine. Three checks assess whether the explanations are stable. The SHAP ranking agrees with permutation importance at a Kendall tau of 0.66. The top-5 feature set is identical across 1,000 bootstrap resamples, with a Jaccard of 1.0. Neutralizing the top-k features drops the mean score monotonically from −0.079 at $k = 1$ to −0.412 at $k = 10$, which confirms that the top features drive the predictions.

Table 11 reports the perturbation study. Content edits move the score in proportion to the semantic weight of the edit. Replacing an entity changes the score by 0.360 on average, a number by 0.345, a date by 0.279, and removing the supporting context by 0.272. The irrelevant insert and clause shuffle leave the score nearly untouched, at 0.006

and 0.001. The top-1 SHAP feature never flips under any operator, and the Spearman correlation stays between 0.84 and 0.95. The sample sizes are small for dates and clause shuffles, 13 and 14 respectively, because not every answer contains an edit-able element, and the table reports n openly.

Table 11
Perturbation stability of the calibrated score and SHAP ranking.

Operator	n	Δscore	Large rate	Top-1 flip	Spearman
Entity replacement	78	0.360	0.372	0.000	0.840
Number replacement	22	0.345	0.364	0.000	0.852
Date replacement	13	0.279	0.308	0.000	0.870
Support removal	66	0.272	0.288	0.000	0.874
Irrelevant insert	120	0.006	0.008	0.000	0.948
Clause shuffle	14	0.001	0.000	0.000	0.921

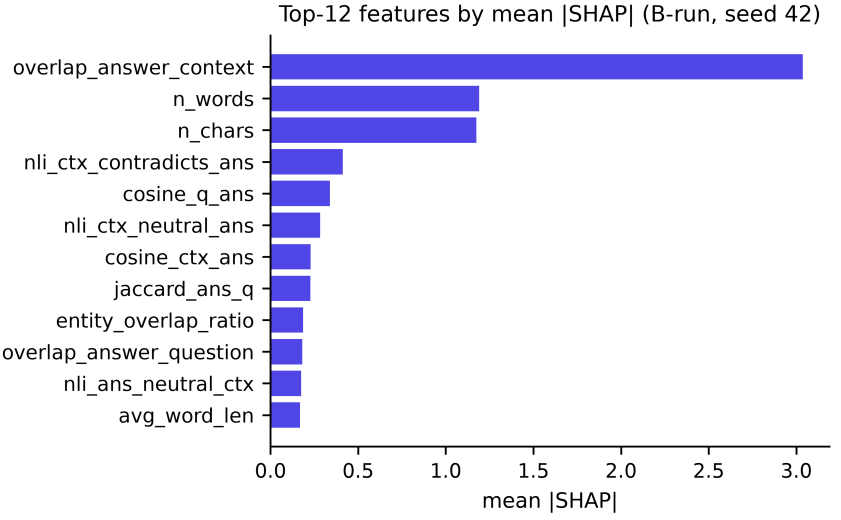


Figure 6: Mean absolute SHAP values for the top-12 features (B-run, seed 42).

A structured expert audit tested these explanations case by case. The B5 run exported 26 error-enriched cases, with 9 false positives, 10 false negatives, and 7 inside the 0.35 to 0.65 band. Two review passes were performed with an AI assistant, which the released sheet records in a review-mode field. Both passes judged the explanation against the answer-context evidence rather than the dataset label, they agreed on 23 of 26 cases (88%), and both marked 15 cases implausible. False positives inherit the length confound, and false negatives occur when a wrong answer reuses context entities, which keeps overlap high. Two audited cases (items 2251 and 2436) carry a hallucinated label even though the context directly supports the answer, so they look like benchmark label noise rather than model errors. The audit sample is deliberately error-heavy, so these proportions describe the audit and not the test set.

5.7. Efficiency and the LLM judge

Table 12 breaks down the latency over 200 test samples. The NLI cross-encoder and the spaCy entity pass dominate at 27.1 ms and 21.1 ms median, while the model itself needs 1.3 ms and SHAP 1.9 ms. The total is 61.8 ms at the median, the deployed artifact weighs 0.82 MB, and the marginal cost of a local prediction is about $0.001 per 1,000 calls. End-to-end HTTP timing gives medians of 42.8 ms for prediction, 57.1 ms for analysis, and 143.6 ms for verification with retrieval over 26 requests, and 502.5 ms with the judge over 6 requests. The component table excludes retrieval, judge, and model-loading time, and the HTTP verification number includes retrieval but not the judge. Local costs assume the reference laptop with amortized hardware and electricity, so they are order-of-magnitude estimates in USD per 1,000

predictions, while the GPT cost comes from the recorded token counts at $0.20 per million input tokens and $1.20 per million output tokens.

Table 12
Per-component latency in ms (n = 200). Components are timed in separate instrumented passes over the same samples, so the component means need not sum to the end-to-end median. Retrieval and judge calls are excluded.

Component	p50	p95	mean
Core lexical	0.10	0.15	0.11
Entity (spaCy NER)	21.13	47.24	25.34
NLI (2 directions)	27.09	68.60	40.09
Semantic (MiniLM)	9.25	15.55	12.56
XGBoost inference	1.26	8.80	2.62
SHAP explanation	1.91	15.52	3.65
Total per sample	61.77		

Table 13 puts every baseline on the same footing. The model-only p50 is timed per row over 200 test rows on the reference laptop, with classical models on CPU and NLI with embeddings on the GPU inside the shared feature pipeline. Classical heads cost almost nothing next to the 61.8 ms feature pipeline, and the two judges are three orders of magnitude slower than any local model. The overlap heuristic needs no model and costs 0.009 ms per row.

Table 13
Per-baseline cost and latency. Size is the seed-42 artifact on disk. Feature models share the 61.8 ms pipeline end to end. The NLI-only head was not persisted, so its cost is the 27.1 ms NLI feature step plus a linear head of the same order as the timed logistic regression. * marks vectorizer-transform-only timings. RoBERTa inference was not timed because the weights were not kept, and its training cost is about 15 minutes per seed. Costs are USD per 1,000 predictions.

Baseline	Inputs	Size	Model p50 (ms)	Cost/1K
Heuristic (overlap)	text	—	0.009	≈$0
Logistic regression	26 feats	1.1 KB	0.001	≈$0.001
Random forest	26 feats	8.4 MB	0.331	≈$0.001
NLI-only	6 NLI probs	—	—	≈$0.001
TF-IDF (answer/full)	text	0.58/4.1 MB	0.014/0.110*	≈$0.001
XGBoost (conventional)	26 feats	0.49 MB	0.005	≈$0.001
RoBERTa-base	256 tokens	125M params	—	training only
EC-XGB (proposed)	35 feats	0.68 MB	0.008	≈$0.001
GPT 5.6 Luna judge	API	—	1443	$0.106
Qwen 3.5 4B judge	4-bit local	4B params	2623	$0

Table 14 compares EC-XGB with two judges on the same 200 balanced HaluEval rows (seed 42). The rows are a stratified subsample of the 3,000-row headline test set. Each class contributes 100 rows, drawn at sampling seed 42, so prevalence is 50%. GPT 5.6 Luna reaches F1 0.846 [0.781, 0.898], Qwen 3.5 4B reaches 0.756 [0.667, 0.824], and EC-XGB reaches 0.990. GPT agrees with EC-XGB on 85.5% of the cases; McNemar is significant ($p = 8.3 \times 10^{-6}$), with 27 judge-wrong/model-right cases against 2 reversed. Qwen agrees with GPT on 88.5% of the rows. An assisted (non-human) review of all 40 Qwen errors finds 20 answers that read as supported under a strict label, 6 ambiguous cases, 12 genuine misses that GPT catches, and 2 false alarms. A length-only baseline reaches F1 0.905 on the same rows, so this comparison does not show semantic detection beyond the length cue. GPT costs $0.106 per 1,000 predictions and needs 1.44 s at the median; Qwen is local and free, and EC-XGB costs about $0.001 per 1,000 and runs in 61.8 ms. The claim is scoped to this sample and these prompts.

5.8. Error analysis

The 3,000-sample test set contains 48 errors, 12 false positives and 36 false negatives, an error rate of 1.6%. A heuristic review of 20 sampled errors (10 false positives and 10 false negatives) assigns them to two categories. Short-answer ambiguity covers very short answers where neither the lexical nor the NLI evidence is decisive, and label ambiguity covers answers that are semantically plausible under both labels. Label ambiguity accounts for 7 false positives and 4 false negatives, and short-answer ambiguity for 3 false positives and 6 false negatives. The sample is

Table 14
Two judges vs. EC-XGB on 200 balanced HaluEval rows (seed 42). GPT uses default decoding and Qwen is local 4-bit. Cost/1K is USD.

Model	Acc	P	R	F1	p50 (ms)	Cost/1K
GPT 5.6 Luna judge	0.865	0.987	0.740	0.846	1443	$0.106
Qwen 3.5 4B judge	0.800	0.969	0.620	0.756	2623	$0
EC-XGB (proposed)	0.990	0.990	0.990	0.990	62	≈$0.001

balanced across the two error types, so the counts describe the reviewed cases rather than the full error population. Both categories trace back to the construction of the benchmark, where short answers are frequently hallucinated by design.

The length distribution behind those errors is the more informative signal. Table 15 bins the full test set by answer length. The hallucinated share is 2.1% among one-word answers and rises to 99.2% above 17 words, so answer length is close to a perfect proxy for the label in HaluEval, and the mean score climbs from 0.024 to 0.993 across the bins. Errors concentrate by length: two-thirds of all false positives fall in the 5-8 word bin, and false negatives concentrate in the 1-4 word bins. This mechanism explains the zero-shot transfer result. Natural RAGTruth responses are full sentences, so the model reads their length as evidence of hallucination, which is why recall reaches 1.0 and AUROC stays near 0.5 on that corpus.

Three checks frame the headline result. A logistic regression on the two length features alone reaches F1 = 0.9457 and AUROC = 0.9759, so a large part of the benchmark signal is answer length. On a length-bin reweighting control with 247 rows from each of the six bins, EC-XGB reaches F1 = 0.9867 and AUROC = 0.9987, and the length-only baseline reaches F1 = 0.9475 and AUROC = 0.9758 on the same 1,482 rows. The difference is +0.039 F1 [0.029, 0.051] and +0.023 AUROC [0.016, 0.031]. Equalizing bin counts does not remove the within-bin association: prevalence still runs from 2.4% in the one-word bin to 99.2% in the 17+ bin, and the positive/negative counts across the six bins are 6/241, 9/238, 33/214, 222/25, 242/5, and 245/2. The bin-equalized comparison is a reweighting control, not proof that the shortcut is gone. A stronger control matches positives and negatives at each exact answer length. On the full test set this gives 163 pairs (326 rows), and the matching is repeated 200 times at different sampling seeds to remove any dependence on one draw. EC-XGB holds an AUROC of 0.9608 with a standard deviation of 0.0022 across those draws, the length-only baseline falls to 0.5257 with a standard deviation of 0.0034, and the gap is 0.4351 [0.429, 0.442]. A length-conditional test removes the cue without subsampling. Inside each exact-length stratum that contains both classes, length is constant, so the length-only model sits at AUROC 0.5 by construction. Across the 12 such strata, EC-XGB reaches a pair-weighted AUROC of 0.9773 with a standard deviation of 0.0009 over seeds.

Table 15
Test-set errors by answer length, mean over seeds (17+ contains only two negative rows).

Words	Rows	Hallucinated	FP rate	FN rate	Mean score
1	524	0.021	0.004	0.182	0.024
2	514	0.033	0.000	0.176	0.035
3–4	474	0.127	0.002	0.178	0.126
5–8	641	0.902	0.106	0.031	0.888
9–16	600	0.982	0.000	0.003	0.980
17+	247	0.992	0.333	0.000	0.993

5.9. Discussion

The benchmark is close to saturation, so the informative result is the behavior outside it. EC-XGB keeps a usable ranking on RAGTruth where the conventional model flags almost every answer, and the display calibrator makes the target scores readable. The gain comes from evidence-consistent features and multi-source training, not from a new learner. A fine-tuned encoder still wins in-domain, so the feature model trades some accuracy for cost and transparency.

6. Conclusion

Hallucination detection needs to be cheap, fast, and explainable before it is useful in practice. This study presented EC-XGB, an evidence-consistent XGBoost risk predictor. It combines 26 evidence-consistency features with eight claim-level NLI aggregates, log-scaled length features under monotone constraints, multi-source training, calibrated probabilities, and SHAP explanations. On HaluEval QA it reaches an F1 of 0.9855 and an AUROC of 0.9984, on par with the conventional XGBoost on a saturated benchmark. The gain appears on the RAGTruth task types used in training, not on unseen tasks. The conventional model flags 99.9% of natural RAGTruth answers as risky and EC-XGB flags 61.3%, with the AUROC rising from 0.475 to 0.582 and the ECE falling from 0.635 to 0.278. A seed-42 display calibrator cuts the RAGTruth QA ECE from 0.726 to 0.126 and needs a labeled target sample (5,034 rows here), yet 252 labeled rows already bring the test ECE to 0.133. One analysis takes 61.8 ms and costs about $0.001 per 1,000 predictions, roughly 100 times less than the GPT 5.6 Luna judge.

The negative results matter as much as the positive ones. A source-only fine-tuned encoder wins the benchmark (F1 0.9912) but ranks worse outside it (RAGTruth AUROC 0.337 against 0.582). The raw model needs no recalibration on HaluEval, the entity and hedging features add no measurable value, and the held-out RAGTruth QA task stays near an F1 of 0.30. The paper therefore frames the supervised model as a domain-specific risk estimator and adaptation check, not a general-purpose detector, and the deployed system adds claim-level verification with evidence retrieval. Future work targets multicalibration under stronger shift, conformal prediction, per-domain threshold transfer, and shortcut-controlled evaluation on unseen tasks.

CRediT authorship contribution statement

Md. Atikur Rahaman: Conceptualization, Data curation, Formal analysis, Investigation, Methodology, Project administration, Resources, Software, Validation, Visualization, Writing - original draft, Writing - review & editing. **Saiful Alam Sabbir:** Writing - review & editing. **MD. Miraz Ahamed:** Writing - review & editing. **Ohidujjaman Tuhin:** Supervision, Validation, Writing - review & editing.

Declarations

- **Funding:** The authors received no specific funding for this work.
- **Conflicts of Interest:** The authors declare no conflicts of interest.
- **Data availability:** The raw evaluation corpora are not redistributed. The repository ships the download scripts, source revisions, SHA-256 hashes, and license notes for HaluEval (MIT), RAGTruth (MIT), and FaithBench (CC BY-NC-SA 4.0). Every split and feature matrix can be regenerated from the original sources. Model artifacts, split indices, and the frozen experiment manifest are included with the code release.
- **Code availability:** The training pipeline, FastAPI inference service, and Next.js dashboard are available in the project repository.
- **Ethical approval:** The study uses publicly released benchmark datasets under their original licenses. It involves no human participants and no personal data, so ethical approval is not required.

Acknowledgements

The authors thank their supervisor, Ohidujjaman Tuhin, for guidance and feedback throughout the study, and the maintainers of HaluEval, RAGTruth, and FaithBench for releasing their datasets under open licenses.

References

[1] P. Manakul, A. Liusie, M. J. Gales, Selfcheckgpt: Zero-resource black-box hallucination detection for generative large language models, in: Proceedings of EMNLP, 2023, pp. 9004–9017. doi:10.18653/v1/2023.emnlp-main.557.

[2] M. Belyi, R. Friel, S. Shao, A. Sanyal, Luna: A lightweight evaluation model to catch language model hallucinations with high accuracy and low cost, in: Proceedings of COLING (Industry Track), 2025, pp. 398–409.
URL https://aclanthology.org/2025.coling-industry.34/

[3] S. Farquhar, J. Kossen, L. Kuhn, Y. Gal, Detecting hallucinations in large language models using semantic entropy, Nature 630 (8017) (2024) 625–630. doi:10.1038/s41586-024-07421-0.

[4] S. Yadav, N. K. Verma, A hybrid framework for hallucination detection in large language models, IEEE Transactions on Artificial Intelligence (2026). doi:10.1109/TAI.2026.3653354.

[5] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Bang, A. Madotto, P. Fung, Survey of hallucination in natural language generation, ACM Computing Surveys 55 (12) (2023) 1–38. doi:10.1145/3571730.

[6] L. Huang, W. Yu, W. Ma, W. Zhong, Z. Feng, H. Wang, Q. Chen, W. Peng, X. Feng, B. Qin, T. Liu, A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions, ACM Transactions on Information Systems 43 (2) (2025) 1–55. doi:10.1145/3703155.

[7] A. Alansari, H. Luqman, Large language models hallucination: A comprehensive survey, Computer Science Review 61 (2026) 100970. doi:10.1016/j.cosrev.2026.100970.

[8] X. Cheng, J. Li, W. X. Zhao, H. Zhang, F. Zhang, F. Zhang, K. Gai, J.-R. Wen, Small agent can also rock! empowering small language models as hallucination detector, in: Proceedings of EMNLP, 2024, pp. 14600–14615. doi:10.18653/v1/2024.emnlp-main.809.

[9] S. Valentin, J. Fu, G. Detommaso, S. Xu, G. Zappella, B. Wang, Cost-effective hallucination detection for llms, arXiv preprint arXiv:2407.21424 (2024).

[10] D. Sundaragiri, L. M. Reddy, P. Gunavardhan, B. Navahith, Framework for hallucination detection in large language models, International Journal of Engineering Research & Technology (IJERT) 15 (04) (2026). doi:10.5281/zenodo.20025987.

[11] I. Haq, M. Saqib, Y. Zhang, I. A. Khan, Quantifying factual divergence in generative models: Shap-lime based hallucination score for llms, Multimedia Systems 32 (2) (2026). doi:10.1007/s00530-025-02150-4.

[12] Y. Deng, Z. Fang, S. Li, L. Chen, Beyond in-domain detection: Spikescore for cross-domain hallucination detection, in: Proceedings of ICLR, 2026.

[13] F. S. Bao, M. Li, R. Qu, et al., Faithbench: A diverse hallucination benchmark for summarization by modern llms, in: Proceedings of NAACL (Short Papers), 2025, pp. 448–461. doi:10.18653/v1/2025.naacl-short.38.

[14] J. Li, X. Cheng, W. X. Zhao, J.-Y. Nie, J.-R. Wen, Halueval: A large-scale hallucination evaluation benchmark for large language models, in: Proceedings of EMNLP, 2023. doi:10.18653/v1/2023.emnlp-main.397.

[15] C. Niu, Y. Wu, J. Zhu, S. Xu, K. Shum, R. Zhong, J. Song, T. Zhang, Ragtruth: A hallucination corpus for developing trustworthy retrieval-augmented language models, in: Proceedings of ACL, 2024, pp. 10862–10878. doi:10.18653/v1/2024.acl-long.585.

[16] S. Lin, J. Hilton, O. Evans, Truthfulqa: Measuring how models mimic human falsehoods, in: Proceedings of ACL, 2022, pp. 3214–3252. doi:10.18653/v1/2022.acl-long.229.

[17] S. Min, K. Krishna, X. Lyu, M. Lewis, W.-t. Yih, P. W. Koh, M. Iyyer, L. Zettlemoyer, H. Hajishirzi, Factscore: Fine-grained atomic evaluation of factual precision in long form text generation, in: Proceedings of EMNLP, 2023, pp. 12076–12100. doi:10.18653/v1/2023.emnlp-main.741.

[18] C. Guo, G. Pleiss, Y. Sun, K. Q. Weinberger, On calibration of modern neural networks, in: Proceedings of ICML, 2017.
URL http://proceedings.mlr.press/v70/guo17a.html

[19] S. M. Lundberg, S.-I. Lee, A unified approach to interpreting model predictions, in: Advances in Neural Information Processing Systems (NeurIPS), 2017.
URL https://papers.nips.cc/paper_files/paper/2017/hash/8a20a8621978632d76c43dfd28b67767-Abstract.html

[20] T. Chen, C. Guestrin, Xgboost: A scalable tree boosting system, in: Proceedings of the 22nd ACM SIGKDD, 2016. doi:10.1145/2939672.2939785.

[21] Y. Bang, Z. Ji, A. Schelten, A. Hartshorn, T. Fowler, C. Zhang, N. Cancedda, P. Fung, HalluLens: LLM hallucination benchmark, in: Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Association for Computational Linguistics, Vienna, Austria, 2025, pp. 24128–24156. doi:10.18653/v1/2025.acl-long.1176.

[22] P. He, J. Gao, W. Chen, DeBERTaV3: Improving DeBERTa using ELECTRA-style pre-training with gradient-disentangled embedding sharing, in: Proceedings of the International Conference on Learning Representations (ICLR), 2023.

[23] S. R. Bowman, G. Angeli, C. Potts, C. D. Manning, A large annotated corpus for learning natural language inference, in: Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2015, pp. 632–642. doi:10.18653/v1/D15-1075.

[24] A. Williams, N. Nangia, S. R. Bowman, A broad-coverage challenge corpus for sentence understanding through inference, in: Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT), 2018, pp. 1112–1122. doi:10.18653/v1/N18-1101.

[25] N. Reimers, I. Gurevych, Sentence-bert: Sentence embeddings using siamese bert-networks, in: Proceedings of EMNLP-IJCNLP, 2019, pp. 3982–3992. doi:10.18653/v1/D19-1410.

[26] W. Wang, F. Wei, L. Dong, H. Bao, N. Yang, M. Zhou, MiniLM: Deep self-attention distillation for task-agnostic compression of pre-trained transformers, in: Advances in Neural Information Processing Systems (NeurIPS), 2020.

[27] M. Honnibal, I. Montani, S. Van Landeghem, A. Boyd, spaCy: Industrial-strength natural language processing in Python (2020). doi:10.5281/zenodo.1212303.

[28] J. C. Platt, Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods, in: A. J. Smola, P. Bartlett, B. Schölkopf, D. Schuurmans (Eds.), Advances in Large Margin Classifiers, MIT Press, Cambridge, MA, 1999, pp. 61–74.

[29] B. Zadrozny, C. Elkan, Transforming classifier scores into accurate multiclass probability estimates, in: Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD), 2002, pp. 694–699. doi:10.1145/775047.775151.

[30] J. Adebayo, J. Gilmer, M. Muelly, I. Goodfellow, M. Hardt, B. Kim, Sanity checks for saliency maps, in: Advances in Neural Information Processing Systems (NeurIPS), 2018.

[31] C. Chen, K. Liu, Z. Chen, Y. Gu, Y. Wu, M. Tao, Z. Fu, J. Ye, INSIDE: LLMs' internal states retain the power of hallucination detection, in: International Conference on Learning Representations (ICLR), 2024.

[32] Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, et al., RoBERTa: A robustly optimized BERT pretraining approach, arXiv preprint arXiv:1907.11692 (2019). doi:10.48550/arXiv.1907.11692.

[33] Y. Ovadia, E. Fertig, J. Ren, Z. Nado, D. Sculley, S. Nowozin, J. V. Dillon, B. Lakshminarayanan, J. Snoek, Can you trust your model's uncertainty? evaluating predictive uncertainty under dataset shift, in: Advances in Neural Information Processing Systems (NeurIPS), 2019.

[34] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, I. Stoica, Judging LLM-as-a-judge with MT-Bench and Chatbot Arena, in: Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track, 2023.